

WineQuality Data Prediction

April 19, 2025

Importing the dependencies

```
[1]: import pandas as pd
```

```
[2]: import numpy as np
```

```
[3]: import matplotlib.pyplot as plt
```

```
[4]: import seaborn as sns
```

```
[5]: from sklearn.model_selection import train_test_split
```

```
[6]: from sklearn.ensemble import RandomForestClassifier
```

```
[7]: from sklearn.metrics import accuracy_score
```

Data Collection

```
[8]: wine_dataset=pd.read_csv('winequality-red.csv')
```

```
[9]: wine_dataset.head()
```

```
[9]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality
0	9.4	5
1	9.8	5

2	9.8	5
3	9.8	6
4	9.4	5

```
[10]: wine_dataset.shape
```

```
[10]: (1599, 12)
```

```
[11]: wine_dataset.isnull().sum()
```

```
[11]: fixed acidity          0
      volatile acidity      0
      citric acid           0
      residual sugar        0
      chlorides             0
      free sulfur dioxide    0
      total sulfur dioxide   0
      density               0
      pH                   0
      sulphates             0
      alcohol               0
      quality               0
      dtype: int64
```

```
[12]: wine_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   fixed acidity         1599 non-null   float64
 1   volatile acidity      1599 non-null   float64
 2   citric acid           1599 non-null   float64
 3   residual sugar        1599 non-null   float64
 4   chlorides             1599 non-null   float64
 5   free sulfur dioxide    1599 non-null   float64
 6   total sulfur dioxide   1599 non-null   float64
 7   density               1599 non-null   float64
 8   pH                   1599 non-null   float64
 9   sulphates             1599 non-null   float64
10   alcohol               1599 non-null   float64
11   quality               1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

```
[13]: wine_dataset.describe()
```

```
[13]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	\
count	1599.000000	1599.000000	1599.000000	1599.000000	
mean	8.319637	0.527821	0.270976	2.538806	
std	1.741096	0.179060	0.194801	1.409928	
min	4.600000	0.120000	0.000000	0.900000	
25%	7.100000	0.390000	0.090000	1.900000	
50%	7.900000	0.520000	0.260000	2.200000	
75%	9.200000	0.640000	0.420000	2.600000	
max	15.900000	1.580000	1.000000	15.500000	

	chlorides	free sulfur dioxide	total sulfur dioxide	density	\
count	1599.000000	1599.000000	1599.000000	1599.000000	
mean	0.087467	15.874922	46.467792	0.996747	
std	0.047065	10.460157	32.895324	0.001887	
min	0.012000	1.000000	6.000000	0.990070	
25%	0.070000	7.000000	22.000000	0.995600	
50%	0.079000	14.000000	38.000000	0.996750	
75%	0.090000	21.000000	62.000000	0.997835	
max	0.611000	72.000000	289.000000	1.003690	

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

```
[14]: wine_dataset.groupby('quality').mean()
```

```
[14]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	\
quality					
3	8.360000	0.884500	0.171000	2.635000	
4	7.779245	0.693962	0.174151	2.694340	
5	8.167254	0.577041	0.243686	2.528855	
6	8.347179	0.497484	0.273824	2.477194	
7	8.872362	0.403920	0.375176	2.720603	
8	8.566667	0.423333	0.391111	2.577778	

	chlorides	free sulfur dioxide	total sulfur dioxide	density	\
quality					
3	0.122500	11.000000	24.900000	0.997464	
4	0.090679	12.264151	36.245283	0.996542	
5	0.092736	16.983847	56.513950	0.997104	
6	0.084956	15.711599	40.869906	0.996615	


7	0.076588	14.045226	35.020101	0.996104
8	0.068444	13.277778	33.444444	0.995212

	pH	sulphates	alcohol
quality			
3	3.398000	0.570000	9.955000
4	3.381509	0.596415	10.265094
5	3.304949	0.620969	9.899706
6	3.318072	0.675329	10.629519
7	3.290754	0.741256	11.465913
8	3.267222	0.767778	12.094444

Visualization


```
[15]: sns.catplot(x='quality',data=wine_dataset,kind='count')
```

```
[15]: <seaborn.axisgrid.FacetGrid at 0x23ec4576810>
```



```
[16]: sns.countplot(x='fixed acidity',hue='quality',data=wine_dataset)
```


```
[16]: <Axes: xlabel='fixed acidity', ylabel='count'>
```



Volatile acidity vs quality

```
[17]: plot=plt.figure(figsize=(5,5))
sns.barplot(x='quality',y='volatile acidity',data=wine_dataset)
```


```
[17]: <Axes: xlabel='quality', ylabel='volatile acidity'>
```



volatile acidity and quality are inversely proportional to each other

```
[19]: plot=plt.figure(figsize=(10,10))  
      sns.barplot(x='quality',y='citric acid',data=wine_dataset)
```

```
[19]: <Axes: xlabel='quality', ylabel='citric acid'>
```



Citric acid is directly proportional to the quality

```
[21]: correlation=wine_dataset.corr()
```

```
[22]: print(correlation)
```

	fixed acidity	volatile acidity	citric acid \
fixed acidity	1.000000	-0.256131	0.671703
volatile acidity	-0.256131	1.000000	-0.552496
citric acid	0.671703	-0.552496	1.000000
residual sugar	0.114777	0.001918	0.143577
chlorides	0.093705	0.061298	0.203823
free sulfur dioxide	-0.153794	-0.010504	-0.060978
total sulfur dioxide	-0.113181	0.076470	0.035533

density	0.668047	0.022026	0.364947
pH	-0.682978	0.234937	-0.541904
sulphates	0.183006	-0.260987	0.312770
alcohol	-0.061668	-0.202288	0.109903
quality	0.124052	-0.390558	0.226373

	residual sugar	chlorides	free sulfur dioxide \
fixed acidity	0.114777	0.093705	-0.153794
volatile acidity	0.001918	0.061298	-0.010504
citric acid	0.143577	0.203823	-0.060978
residual sugar	1.000000	0.055610	0.187049
chlorides	0.055610	1.000000	0.005562
free sulfur dioxide	0.187049	0.005562	1.000000
total sulfur dioxide	0.203028	0.047400	0.667666
density	0.355283	0.200632	-0.021946
pH	-0.085652	-0.265026	0.070377
sulphates	0.005527	0.371260	0.051658
alcohol	0.042075	-0.221141	-0.069408
quality	0.013732	-0.128907	-0.050656


	total sulfur dioxide	density	pH	sulphates \
fixed acidity	-0.113181	0.668047	-0.682978	0.183006
volatile acidity	0.076470	0.022026	0.234937	-0.260987
citric acid	0.035533	0.364947	-0.541904	0.312770
residual sugar	0.203028	0.355283	-0.085652	0.005527
chlorides	0.047400	0.200632	-0.265026	0.371260
free sulfur dioxide	0.667666	-0.021946	0.070377	0.051658
total sulfur dioxide	1.000000	0.071269	-0.066495	0.042947
density	0.071269	1.000000	-0.341699	0.148506
pH	-0.066495	-0.341699	1.000000	-0.196648
sulphates	0.042947	0.148506	-0.196648	1.000000
alcohol	-0.205654	-0.496180	0.205633	0.093595
quality	-0.185100	-0.174919	-0.057731	0.251397

	alcohol	quality
fixed acidity	-0.061668	0.124052
volatile acidity	-0.202288	-0.390558
citric acid	0.109903	0.226373
residual sugar	0.042075	0.013732
chlorides	-0.221141	-0.128907
free sulfur dioxide	-0.069408	-0.050656
total sulfur dioxide	-0.205654	-0.185100
density	-0.496180	-0.174919
pH	0.205633	-0.057731
sulphates	0.093595	0.251397
alcohol	1.000000	0.476166
quality	0.476166	1.000000

Constructing heat map to understand the correlation between the columns

```
[24]: plt.figure(figsize=(10,10))
sns.heatmap(correlation,cbar=True,square=True,fmt='.
↪1f',annot=True,annot_kws={'size':8},cmap='Blues')
```

[24]: <Axes: >



Data preprocessing

```
[25]: X=wine_dataset.drop(['quality'],axis=1)
```

```
[27]: print(X)
```

```

      fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0              7.4              0.700              0.00              1.9              0.076
```

1	7.8	0.880	0.00	2.6	0.098
2	7.8	0.760	0.04	2.3	0.092
3	11.2	0.280	0.56	1.9	0.075
4	7.4	0.700	0.00	1.9	0.076
...
1594	6.2	0.600	0.08	2.0	0.090
1595	5.9	0.550	0.10	2.2	0.062
1596	6.3	0.510	0.13	2.3	0.076
1597	5.9	0.645	0.12	2.0	0.075
1598	6.0	0.310	0.47	3.6	0.067

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
0	11.0	34.0	0.99780	3.51	0.56
1	25.0	67.0	0.99680	3.20	0.68
2	15.0	54.0	0.99700	3.26	0.65
3	17.0	60.0	0.99800	3.16	0.58
4	11.0	34.0	0.99780	3.51	0.56
...
1594	32.0	44.0	0.99490	3.45	0.58
1595	39.0	51.0	0.99512	3.52	0.76
1596	29.0	40.0	0.99574	3.42	0.75
1597	32.0	44.0	0.99547	3.57	0.71
1598	18.0	42.0	0.99549	3.39	0.66

	alcohol
0	9.4
1	9.8
2	9.8
3	9.8
4	9.4
...	...
1594	10.5
1595	11.2
1596	11.0
1597	10.2
1598	11.0

[1599 rows x 11 columns]

Label binarisation

```
[31]: Y=wine_dataset['quality'].apply(lambda y_value: 1 if y_value>=7 else 0)
```

```
[32]: print(Y)
```

0	0
1	0
2	0
3	0

```
4      0
      ..
1594   0
1595   0
1596   0
1597   0
1598   0
```

Name: quality, Length: 1599, dtype: int64

Splitting the data into training and test data

```
[33]: train_x, test_x, train_y, test_y = train_test_split(X, Y, test_size=0.
      ↪ 2, stratify=Y, random_state=2)
```

```
[34]: print(X.shape, train_x.shape, train_y.shape, test_x.shape, test_y.shape)

(1599, 11) (1279, 11) (1279,) (320, 11) (320,)
```

Model Training using RandomForest Classifier Model

```
[35]: model = RandomForestClassifier()
```

```
[36]: model.fit(train_x, train_y)
```

```
[36]: RandomForestClassifier()
```

```
[37]: train_x_prediction = model.predict(train_x)
```

```
[38]: train_x_accuracy = accuracy_score(train_x_prediction, train_y)
```

```
[39]: print(train_x_accuracy)
```

1.0

```
[40]: test_x_prediction = model.predict(test_x)
```

```
[41]: test_x_accuracy = accuracy_score(test_x_prediction, test_y)
```

```
[42]: print(test_x_accuracy)
```

0.9375

```
[45]: X_new = test_x.iloc[2]
```

```
[47]: nparray = np.asarray(X_new)
```

```
[48]: reshaped = nparray.reshape(1, -1)
```

```
[49]: X_new_df = pd.DataFrame(reshaped, columns=train_x.columns)
```

```
[50]: prediction = model.predict(X_new_df)
```

```
[51]: print(prediction)
```

```
[0]
```

```
[52]: print(test_y.iloc[2])
```

```
0
```

```
[60]: input_data=(7.5,0.5,0.36,6.1,0.071,17.0,102.0,0.9978,3.35,0.8,10.5)
```

```
[61]: nparray=np.asarray(input_data)
```

```
[62]: reshaped=nparray.reshape(1,-1)
```

```
[63]: input_data_df=pd.DataFrame(reshaped,columns=train_x.columns)
```

```
[64]: prediction=model.predict(input_data_df)
```

```
[65]: print(prediction)
```

```
[0]
```

```
[66]: if (prediction[0]==1):  
        print('Good Quality Wine')  
    else:  
        print('Bad Quality Wine')
```

```
Bad Quality Wine
```

```
[ ]:
```